1 Linear Regression

In a regression problem we have predictor variables $a_1, \ldots, a_d$ and a measured variable b. In linear regression, we assume there is a relation $b \approx \sum_i a_i x_i$ for some $x_1, \ldots, x_d \in \mathbb{R}$.

We assume we received n batches $(a_{i,1}, \ldots, a_{i,d}, b_i), i = 1..n$. In the least square method we minimize

$$\sum_i (a_{i,1} x_1 + \ldots a_{i,d} x_d - b_i)^2.$$

Formally:

Definition 1. *On the input we have $A = \mathbb{R}^{n \times d}$ and $b \in \mathbb{R}^n$. Least square linear regression asks for $x \in \mathbb{R}^d$ so that*

$$\|Ax - b\|_2$$

is minimized.

1.1 Exact solution

Assume $b = Ax' + b'$ where b' is orthogonal to column space of A and Ax' is projection of b onto column space of A. Then (by Pythagoraean theorem)

$$\|Ax - b\|_2^2 = \|A(x - x') - b'\|_2^2 = \|A(x - x')\|_2^2 + \|b'\|_2^2$$

which is minimized when $x = x'$. The condition of Ax' being projection is equivalent to

$$A^T(Ax' - b) = A^T b' = 0$$

so equivalently we have a following condition

$$A^T A x' = A^T b. \tag{1}$$

If $(A^T A)$ is invertible (its rank is d), we can simply compute

$$x' = (A^T A)^{-1} A^T b.$$

Definition 2. *Let $A = U \Sigma V^T$ be SVD of A. Let $\Sigma^\dagger$ be defined as diagonal matrix where $\Sigma_{i,i}^\dagger = \frac{1}{\Sigma_{i,i}}$ if $\Sigma_{i,i} \neq 0$ and 0 otherwise. We then call $A^\dagger = V \Sigma^\dagger U^T$ a pseudoinverse of A.*

Theorem 3. *$x' = A^\dagger b$ satisfies condition (1) and has minimal L_2 norm among all the solutions.*

Proof. First part:

$$A^T A x' = A^T A A^\dagger b = (V\Sigma^T U^T)(U\Sigma V^T)(V\Sigma^\dagger U^T)b = V\Sigma^T\Sigma\Sigma^\dagger U^T b = V\Sigma^T U^T b = A^T b$$

(note, $\Sigma^T\Sigma\Sigma^\dagger = \Sigma^T$ even though $\Sigma\Sigma^\dagger \neq I$ generally)

Second part: any solution is of the form

$$x'' = A^\dagger b + z$$

where $A^T A z = 0$, or equivalently $V\Sigma^T\Sigma V^T z = 0$ or $\Sigma^T\Sigma V^T z = 0$ (since V is orthonormal) or $\Sigma V^T z = 0$ (since $\text{Ker}(\Sigma^T\Sigma) = \text{Ker}(\Sigma)$) or $V^T z \in \text{Ker}(\Sigma)$ or $z \in V \cdot \text{Ker}(\Sigma)$.

We have $A \dagger b = V\Sigma^\dagger U^T b \in V \cdot \text{Im}(\Sigma^\dagger)$.

Both $\text{Ker}(\Sigma)$ and $\text{Im}(\Sigma^\dagger)$ are subspaces of $\mathbb{R}^d$ - and by closer investigation since Σ is diagonal, both depend on zero/non-zero elements in $\Sigma_{i,i}$. Namely, $\text{Ker}(\Sigma) \perp \text{Im}(\Sigma^\dagger)$, and since V is orthonormal, $V \cdot \text{Ker}(\Sigma) \perp V \cdot \text{Im}(\Sigma^\dagger)$. So by the Pythagorean theorem,

$$\|x''\|_2^2 = \|A^\dagger b\|_2^2 + \|z\|_2^2 \geq \|A^\dagger b\|_2^2$$

which proves optimality. $\qquad\square$

Downside: time to compute SVD is $\mathcal{O}(\min(n^2 d, nd^2))$ which can be prohibitive.

1.2 Approximate solution

Instead of solving exact regression, we pick a projection $\Pi \in \mathbb{R}^{m\times n}$ and solve a problem of smaller dimensionality (m instead of n):

$$\text{minimize} \quad \|\Pi A x - \Pi b\|_2$$

It is enough to use subspace embedding Π for space spanned on columns of A + single vector b. Thus we can pick oblivious subspace embedding for $m = \mathcal{O}(d/\varepsilon^2)$, and have

$$\forall_{x\in\mathbb{R}^d}\|Ax - b\|_2 \leq \|\Pi Ax - \Pi b\|_2 \leq (1+\varepsilon)\|Ax - b\|_2.$$

Thus minimizing projected problem provides $1 + \varepsilon$ approximation to original regression problem.

Total computation time is $\mathcal{O}(mn + \min(m^2 d, md^2)) = \mathcal{O}(nd/\varepsilon^2 + d^3/\varepsilon^2)$.

2 Low rank approximation

Consider input matrix $A \in \mathbb{R}^{n\times d}$. The goal of the low-rank approximation is the following: find B such that B has small rank and $B \approx A$.

Denote such $B = C \times D$, where $C \in \mathbb{R}^{n\times k}$ and $D \in \mathbb{R}^{k\times d}$. Motivation (assume k is small)

- B requires much less space to store: $nk + kd$ vs nd.

- matrix-vector multiplication involving B is much faster: $B \cdot v$ takes $\mathcal{O}(nk + kd)$ time vs $\mathcal{O}(nd)$ time of $A \cdot v$.

- matrix-matrix multiplication: for $X \in \mathbb{R}^{d\times m}$, $B \cdot X$ takes $\mathcal{O}(kdm + nkm)$, vs $\mathcal{O}(ndm)$ time of $A \cdot X$

- A might have structure + noise, B is denoising of A

2.1 Exact solution

We are looking at

$$\arg \min_{B:\text{rank}(B)\leq k} \|A - B\|$$

and denote it as A_k, best rank-k approximation of A.

How to find such A_k? Following theorem holds for both $\|\cdot\|_F$ and $\|\cdot\|_2$ norms.

Theorem 4. *Consider SVD of $A = U\Sigma V^T$. Let Σ_k be Σ where only k largest in absolute value singular values are preserved, and every other value is zeroed.*

$$A_k = U\Sigma_k V^T \tag{2}$$

Proof. TODO $\qquad\square$

Unfortunately the time is dominated by SVD computation $\mathcal{O}(\min(nd^2, n^2d))$.

2.2 Approximate solution

Obtaining good approximate solution is possible for this problem, using the same framework: we project our problem to smaller dimension and hope that solution in reduced dimension approximates good solution to original problem.

Specifically, we use projection matrix $S \in \mathbb{R}^{m\times n}$, for small m. So $m = \mathcal{O}(k/\varepsilon^2)$ (note, m is independent of dimensions of A, and depends on desired rank k).

We mention following theorem

Theorem 5 ([CW13]). *There is algorithm that outputs A'_k of rank k in a factorized form, satisfying $\|A'_k - A\|_F \leq \|A_k - A\|_F \cdot (1 + \varepsilon)$. The time to compute is $\mathcal{O}(nnz(A) + (n + d)poly(k/\varepsilon))$*

We skip the proof due to lack of time/space.

3 Sparse Fourier

Fourier transform: signal $\to$ frequencies.

Definition 6. *Assume $a = (a_0, \ldots, a_{n-1})$ is a signal. Let ω be n-th root of unity, that is $\omega = e^{-\frac{2\pi}{n}}$. Let F be such that $F_{ij} = \frac{1}{\sqrt{n}}\omega^{ij}$. Then $\hat{a} = Fa$ is a (Discrete) Fourier transform of a.*

DFT can be computed in $\mathcal{O}(n\log n)$ time. However, for some applications this time can already be prohibitive. Consider a following scenario (of signal compression):

- Take input signal a and compute $\hat{a}$.

- Let $\hat{a}_k$ be $\hat{a}$ with only k largest magnitude elements kept (rest is zeroed).

- Output $a_k = F^{-1}\hat{a}_k$.

If we consider complexity measure of Fourier support $fs(a) = \|\hat{a}\|_0$ that is number of non-zero Fourier coefficients, then actually there is

$$a_k = \arg \min_{x:fs(x)\leq k} \|a - x\|_2$$

(proof: exercise)

If we assume that a comes from real-life scenarios (photos, audio recording), then it should have only few "strong" frequencies, rest is noise. Since a_k has much simpler representation (namely, $\hat{a}_k$ which takes $\mathcal{O}(k \log n)$ bits), this is a lossy compression scheme.

How do we compute $\hat{a}_k$ efficiently? (Assumption is that we have random access to a, otherwise just *reading* the input would dominate the computation time.)

Simpler question: can we recover $\hat{a}$ if we know that $fs(a) = \|\hat{a}\|_0 \leq k$ (so there are only k non-zero frequencies)?

3.1 Recovering sparse signal

Assume wlog n is power of two.

Define $p_{d,\ell}(x) = \sum_{i:i \bmod 2^\ell = d} \hat{a}_i x^i$, and for a polynomial $p(x)$ we write $\|p\|^2 = \sum_i p_i^2$.

First trick is how to estimate $\|p_{d,\ell}\|_2^2$ with additive error. Given such blackbox, we can proceed:

1. Define $S_\ell = \{i : \|p_{i,\ell}\|^2 > 0\}$

2. Given S_ℓ, compute $S_{\ell+1}$: for each $d \in S_\ell$, test for $e \in \{d, d + 2^\ell\}$ whether $\|p_{e,\ell+1}\|^2 > 0$ and if so, add e to $S_{\ell+1}$.

3. $p_{i,\log n} = \hat{a}_i$

The idea is to start at level 0 and proceed. The size of each S_ℓ is bounded by k, thus the total number of steps (2) done is $\mathcal{O}(k \log n)$.

References

[CW13] Kenneth L. Clarkson and David P. Woodruff. Low rank approximation and regression in input sparsity time. In *Symposium on Theory of Computing Conference, STOC'13, Palo Alto, CA, USA, June 1-4, 2013*, pages 81–90, 2013.